

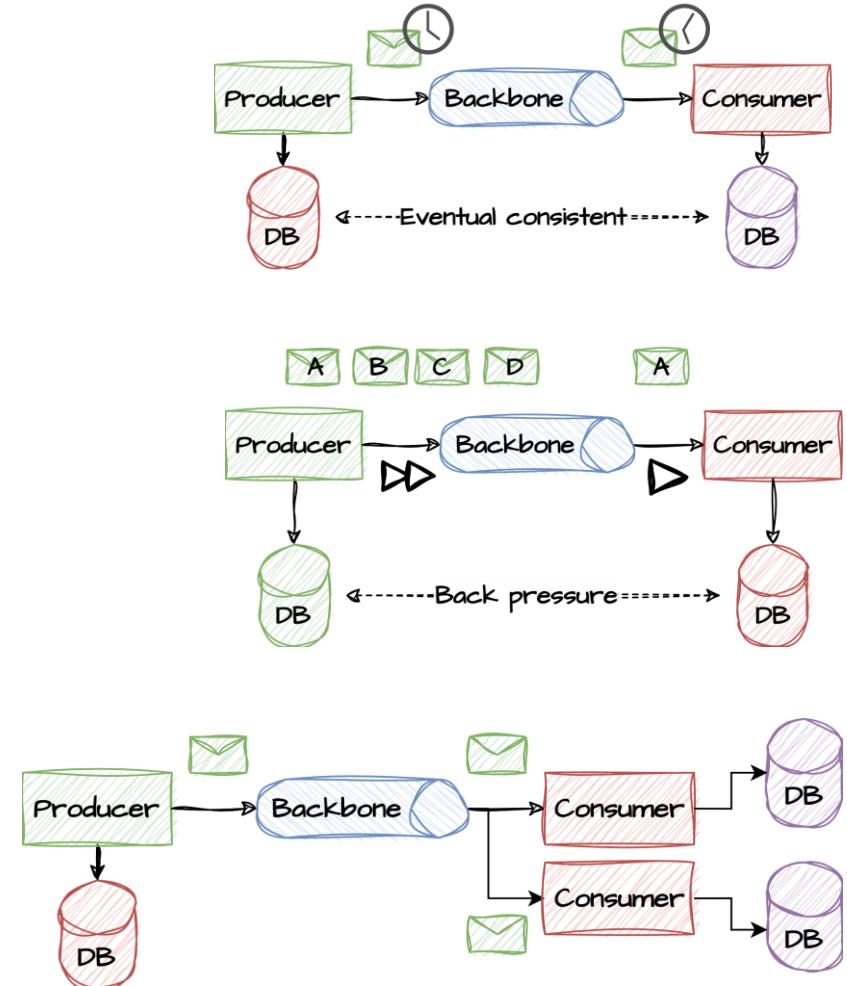
Event driven architecture

EDA overview

Event-driven architecture (EDA) is a software design pattern that enables an organization to detect “events” or important business moments (such as a transaction, site visit etc) and act on them in real time or near real time. This pattern replaces the traditional “request/response” architecture where services would have to wait for a reply before moving on to the next task. The synchronous communication approach does not scale well and introduces blocking code and increases risks of failures.

In EDA, events are emitted by an event emitter and broadcasted to any number of listeners through an event backbone. The event emitter has no knowledge of whom may consume the event. Event listeners are also completely decoupled from the event emitters and may emit new events onto the backbone as a response to the event. The decoupling of producers and consumers adds several benefits to the system design. The decoupling allows producers and consumers to produce/consume at different speed rates, allow new consumers to be added seamlessly and replay historic events.

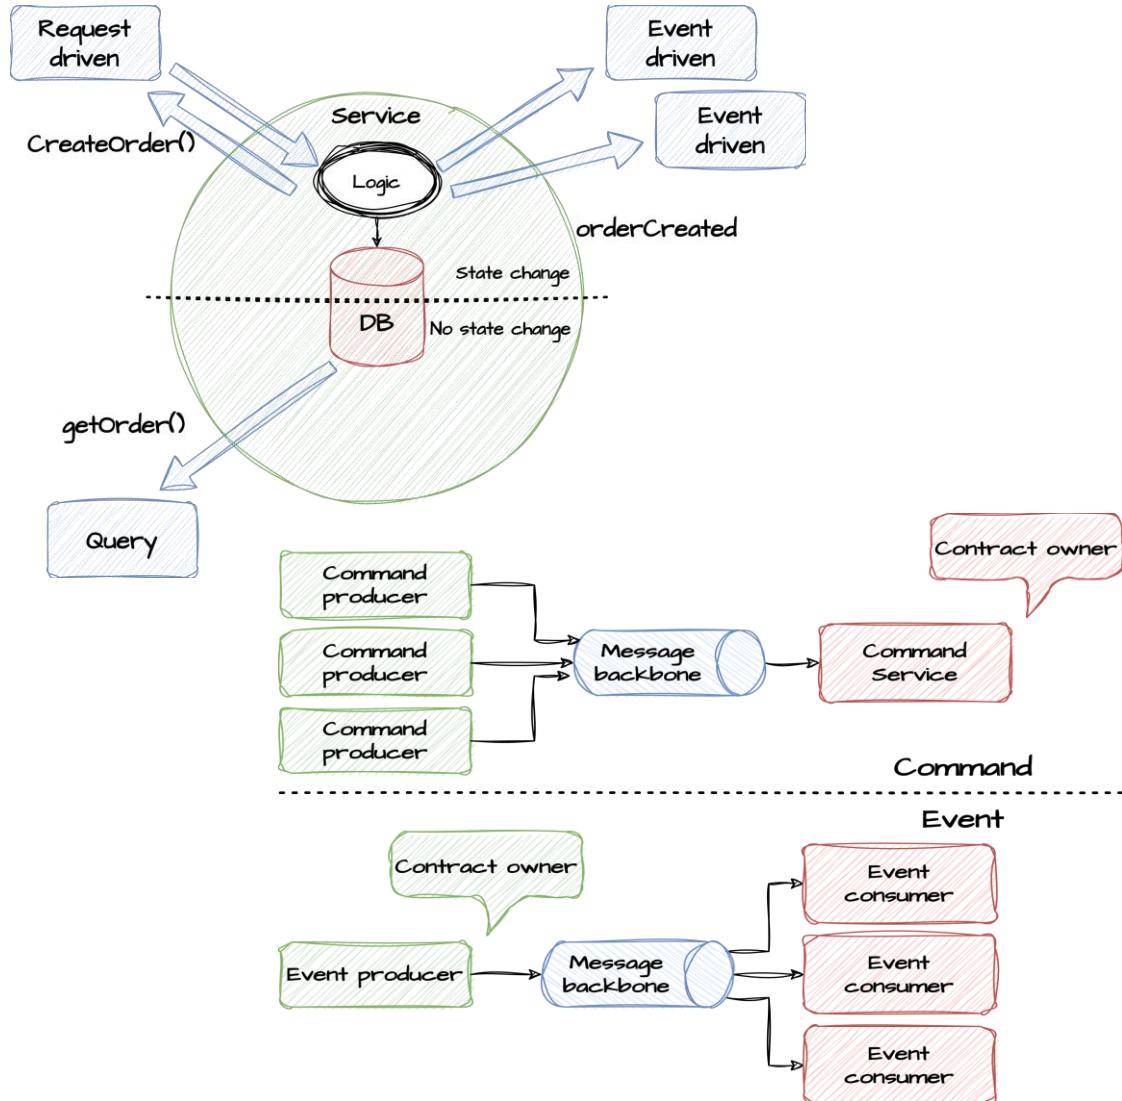
The downside is added complexity when introducing an event pipeline in between components, observability issues and versioning of data.



Events, Commands and Queries

Events

- Something happened (provide identity + context)
- Zero or many consumers (completely decoupled)
- Single producer (owner of the event and its data)
- Owned by producer



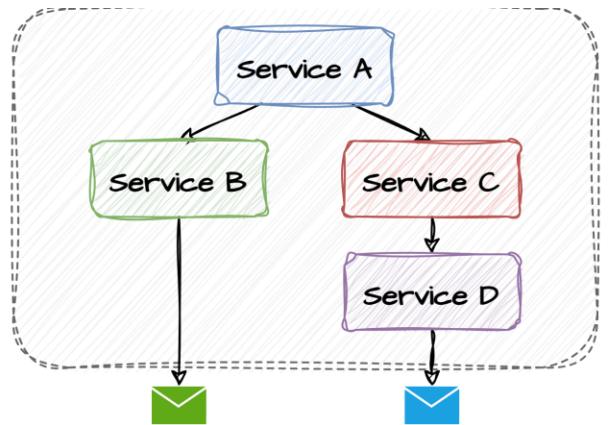
Commands

- Invoke behavior (target a specific recipient)
- Single consumer (consumer executes the command)
- Multiple senders (Anyone can send a command)
- Owned by consumer

Queries

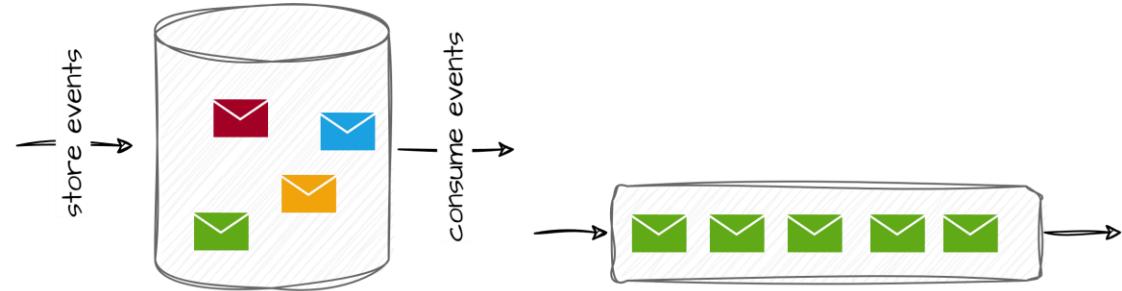
- Peer-to-peer. Typically in the form of HTTP calls
- No side-effects (state changes) in the service.
- Read-only and typically backed by a data-store

Event communication patterns



Process Manager

Central unit that
orchestrates communication



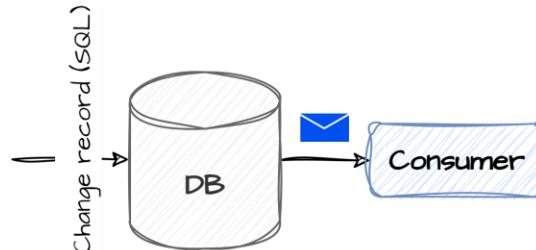
Event Sourcing

Store event and build
current state using events



Point-to-point

Sender put message on queue
Receiver reads from queue and deletes message



Change data capture

Consume changes directly from DB



Pub/Sub

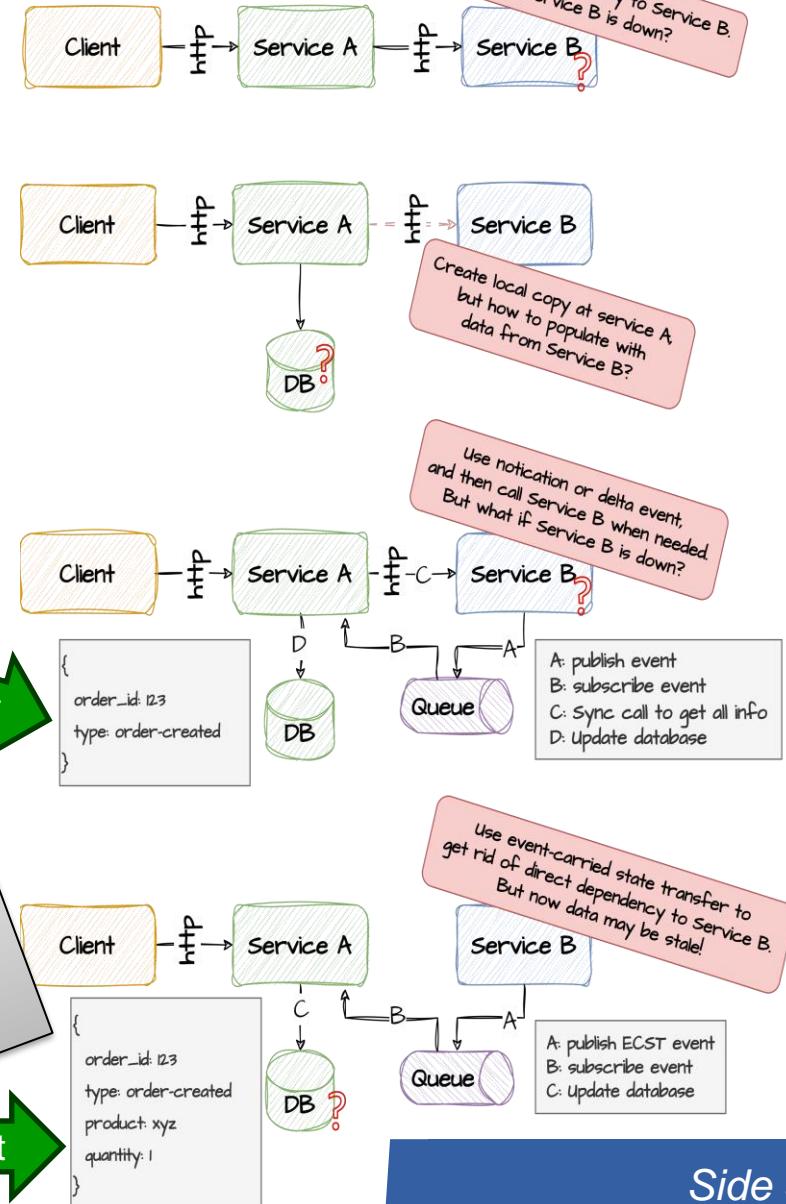
Publish same message to
many consumers. Messages preserved

Events

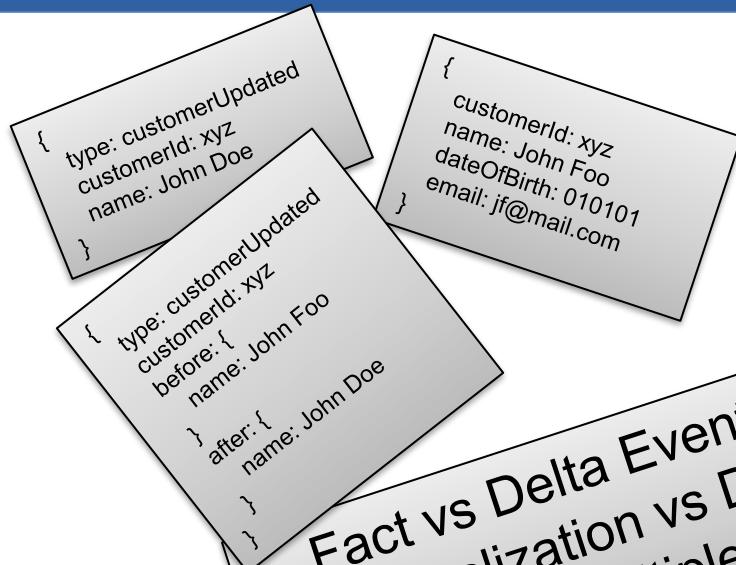


Events is about something that has happened. They are **ordered** and **immutable** and once published cannot be changed or re-ordered. **Multiple consumers** can subscribe to these events. Event may carry more (**fact**) or less (**delta**) information which is a trade-off between high-coupling due to additional lookups (consistency) or risk of carrying stale data but with no additional lookup (availability).

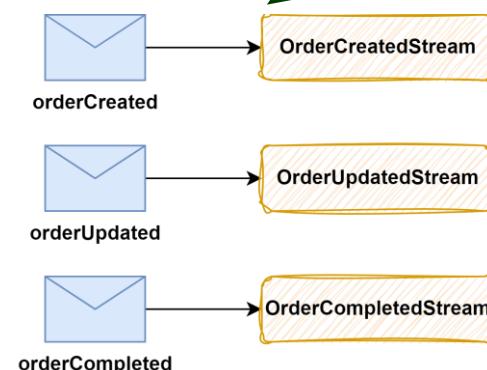
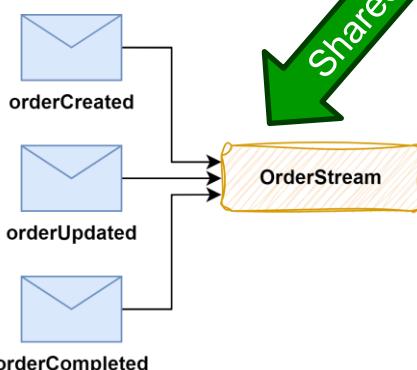
Events should only carry data that it owns. To prevent model sprawl, its the responsibility of the consumer to consume data from other referenced domains.



Designing events

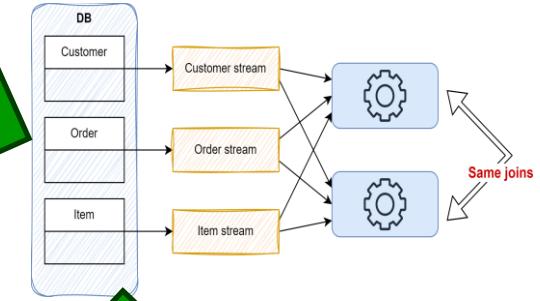


- Fact vs Delta Event types
- Normalization vs Denormalization
- Single vs Multiple Event types pr topic

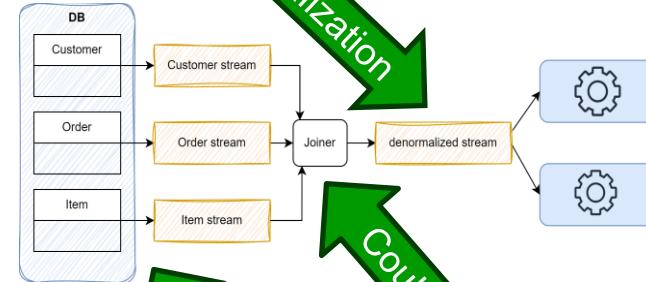


One topic pr event

High coupling to datasource



De-normalization

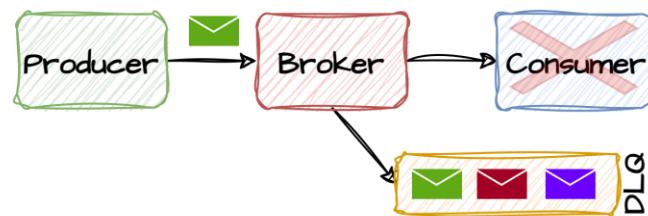


Could also be trans outbox

CDC

Order of event types is lost

Event delivery failures



Dead-letter queues

Event fails to reach target, message can be sent to special queue for retry/inspection



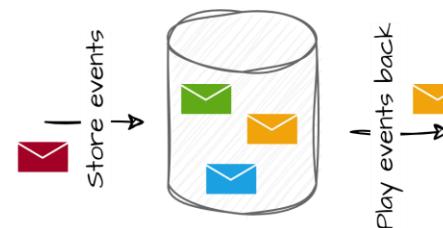
Redrive policies

Sometimes the broker may retry delivery for you based on certain retry criterias



Dropping events

Sometimes it makes sense to just drop events if they cannot reach target system



Archive / replay

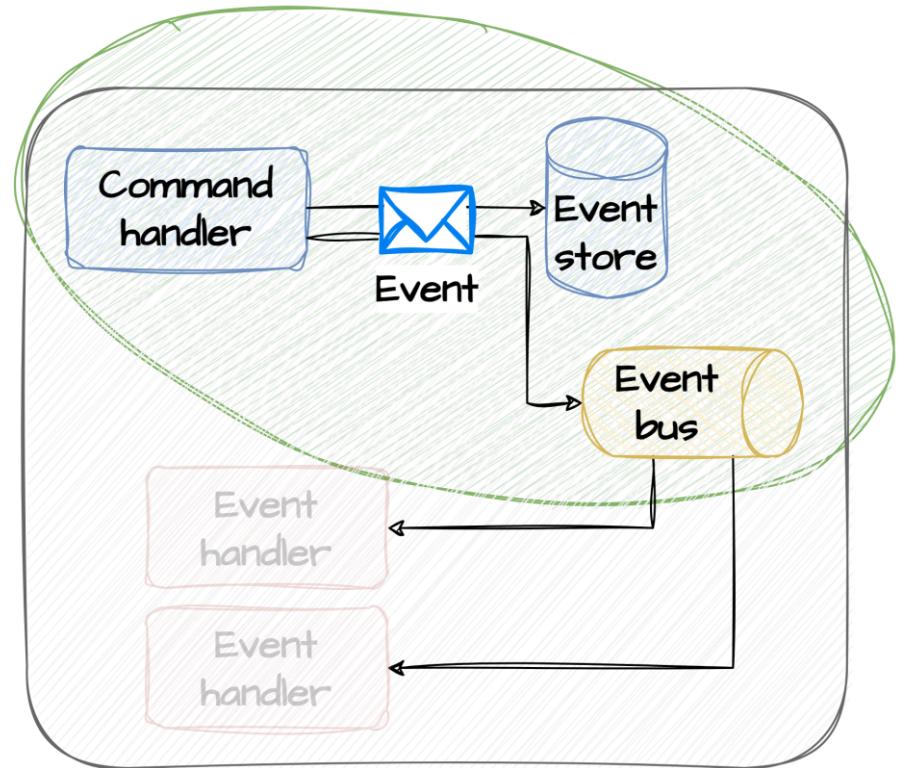
Store events and retry later
Make sure consumers are idempotent

Event Sourcing

Event sourcing is a way of modeling and persisting events (typically delta events) as a series of **immutable** and **ordered** events rather than focusing on the current state of the entity behind the event.

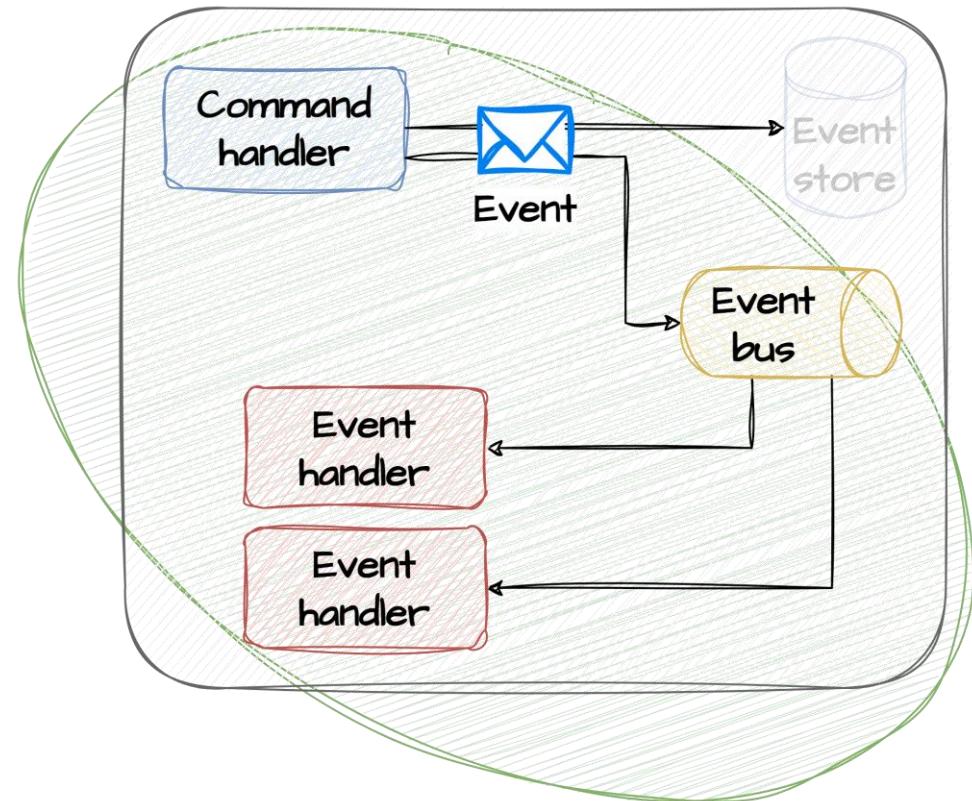
Events are stored in an event store which is local to the service (bounded context) that it lives in. Events in the event store can be replayed and thereby used for state reconstruction or used for auditing or debugging purposes.

To optimize replay and state reconstruction sometimes snapshots are made. When recreating state, you find the latest snapshot and apply all changes made since then.



Event Streaming

Event streaming is an **ordered, unbounded** sequence of **immutable** events flowing from producers to consumers through an event backbone with an **at-least-once** delivery guarantee. Whenever a service or bounded context makes changes to an entity that it owns, this can be published as an event on an event backbone. The event may carry context depending on the design. Event streaming uses a pub-sub approach to enable decoupled communication between systems. Consumers subscribe to a topic/channel on the event pipeline, and producers produce events to these topics/channels. Services may act as both consumers and producers of events. The pub-sub design decouples the producers and consumers, making it easier to scale each part of the system individually.



Event Projection - CQRS

CQRS, which stands for **Command Query Responsibility Segregation**, is an architectural pattern used to separate the concerns of handling commands (write operations) and queries (read operations) within an application.

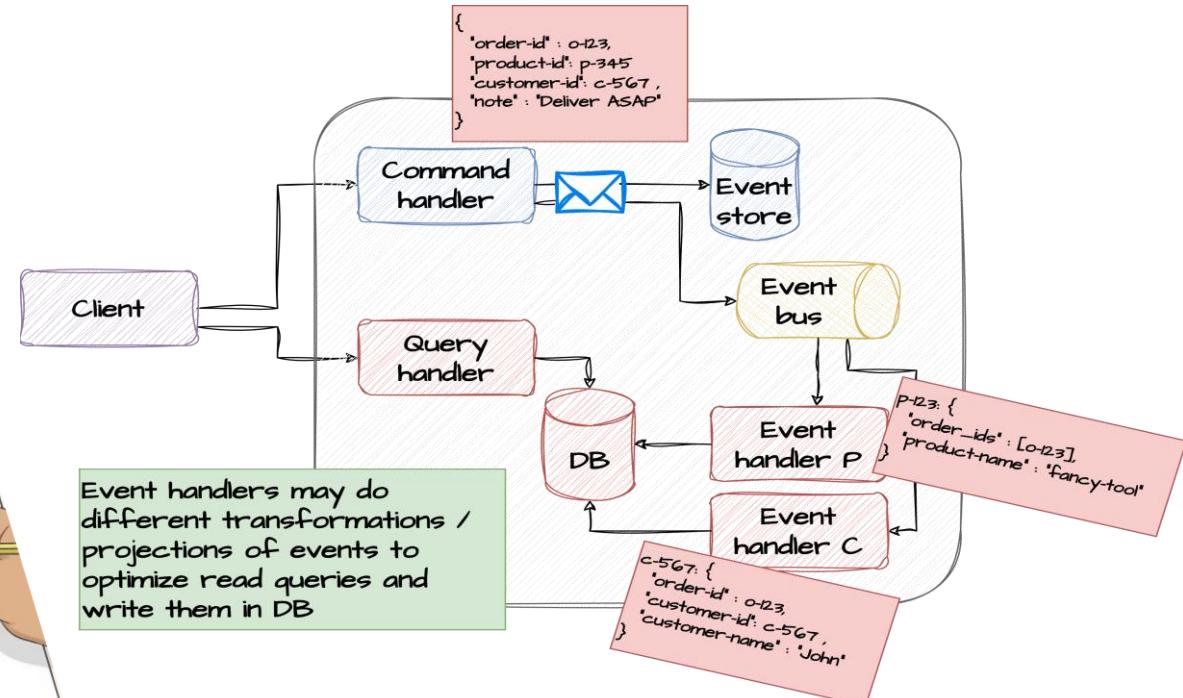
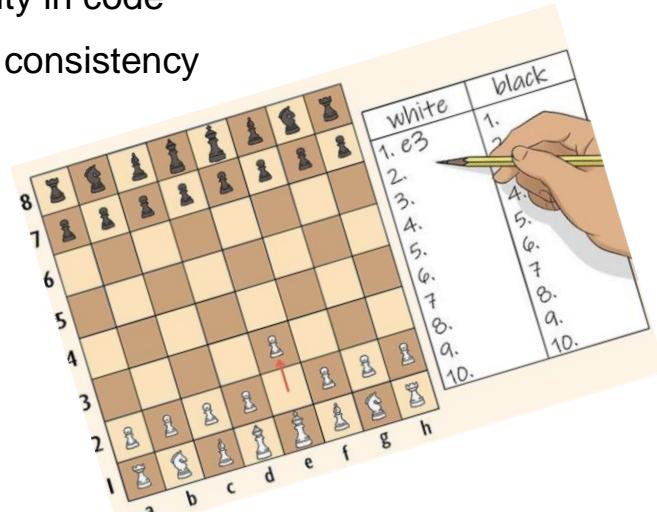
By separating read/write operations this way you can optimize each part independently. Read operations typically have different forms and transformers (event handlers) can project the same write operations into different projections that are read optimized.

Advantages:

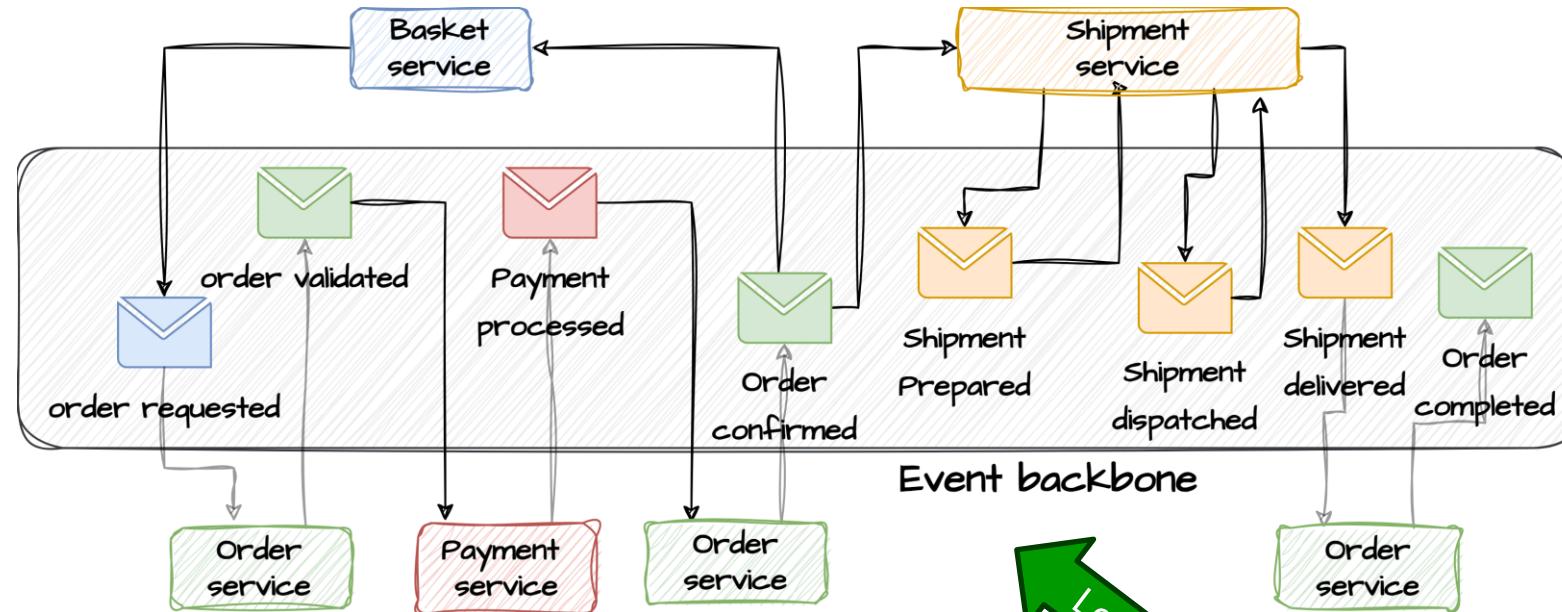
- Scalability – Separate read/write operations
- Flexibility – Create different projections

Disadvantages:

- Complexity in code
- Eventual consistency



EDA at a larger scale



Looks simple – but is it really?

DDD and Bounded contexts

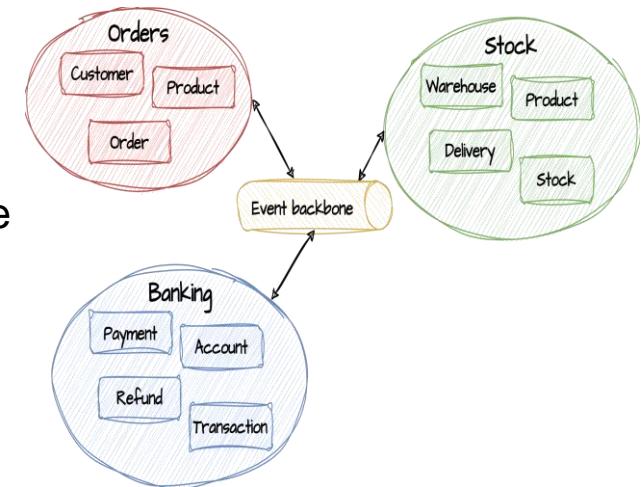
DDD (domain driven development) is about technical- and domain model alignment so that domain experts and technical team members working together can communicate about the model in an unambiguous way.

In DDD, **bounded contexts** help to establish clear boundaries between different parts of the system by grouping related domain **entities**. This allows for better separation of work as smaller teams can focus on sub domains. Inside a boundary, the model must be **consistent** and **coherent**, meaning that the entities must be inter-related and dependent on each other. Messages sent between services inside a bounded context is referred to as **domain events**.

The same entity name may occur within different bounded contexts (like Product in Orders and Stock).

A public language (model) needs to be specified when communicating between bounded contexts.

Events sent between bounded context are referred to as **integration events**. Sometimes translations to/from integration events must happen before sending to the backbone or receiving from the backbone if the producing or consuming bounded context is not aware of integration event language.



Messages between bounded context



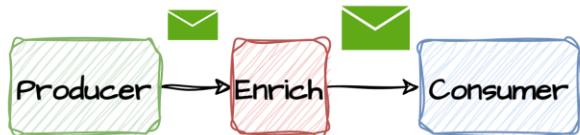
Conformist

Don't translate message, consume as is



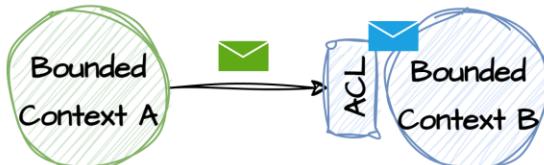
Open-Host Service

Producer does not conform to contract
Messages get translated prior to publishing



Enrich

Sometimes you want consumers getting more info vs producers providing it



Anti-corruption

Consumer does not conform to contract
Messages get translated before consumption